

Modelling Call-Time Choice as Effect using Scoped Free Monads

Niels Bunkenburg

Master's Thesis
Programming Languages and Compiler Construction
Department of Computer Science
Kiel University

Advised by
Priv.-Doz. Dr. Frank Huch
M. Sc. Sandra Dylus

February 11, 2019



Erklärung der Urheberschaft

Ich erkläre hiermit an Eides statt, dass ich die vorliegende Arbeit ohne Hilfe Dritter und ohne Benutzung anderer als der angegebenen Hilfsmittel angefertigt habe. Aus fremden Quellen direkt oder indirekt übernommene Gedanken sind als solche kenntlich gemacht. Die Arbeit wurde bisher in gleicher oder ähnlicher Form in keiner anderen Prüfungsbehörde vorgelegt und auch noch nicht veröffentlicht.

Ort, Datum

Unterschrift

Abstract

Contents

1	Introduction	1
2	Preliminaries	2
2.1	Coq	2
2.2	Haskell	2
2.2.1	Monad and MonadPlus	2
2.3	Curry	2
2.3.1	Non-strictness	2
2.3.2	Sharing	2
2.3.3	Non-determinism	2
2.4	Modelling Curry Programs using Monadic Code Transformation	2
3	Call-Time Choice modelled in Haskell	7
3.1	Free Monads	8
3.2	Modelling Effects	9
3.2.1	Combining Effects	9
3.2.2	Simplified Pattern Matching	11
3.2.3	Handling Effects	13
3.3	Sharing	17
4	Call-Time Choice modelled in Coq	18
4.1	Non-strictly Positive Occurrence	18
4.2	Containers	19
4.3	Modelling Effects	20
4.4	Sharing	20
5	Curry Programs modelled in Coq	21
6	Conclusion	22

1 Introduction

2 Preliminaries

2.1 Coq

- Introduce the necessary Coq concepts to understand the paper

2.2 Haskell

- Introduce the necessary Haskell concepts to understand the paper

2.2.1 Monad and MonadPlus

2.3 Curry

- Introduce the necessary Curry concepts to understand the paper

2.3.1 Non-strictness

2.3.2 Sharing

2.3.3 Non-determinism

2.4 Modelling Curry Programs using Monadic Code Transformation

- Why is the naive MonadPlus approach not sufficient to model Curry semantic?
- Motivate usage of monadic data types
- Introduce explicit sharing

Modelling Curry programs in a language like Haskell requires a transformation of non-deterministic code into a semantically equivalent, deterministic program. First, we have a look at the direct representation of non-determinism used in the KiCS2 implementation as described by Braßel et al. [2011].

Non-determinism in Curry is not limited to flat non-determinism but can occur within components of data structures and anywhere in a computation. This means that expressing non-determinism via Haskell's list monad is not sufficient to model Curry's non-determinism. Instead, existing data types receive additional constructors that represent

TODO: Example

failure and the choice between two values. For example, the extended list data type looks as follows.

```
data List a = Nil | Cons a (List a) | Choice (List a) (List a) | Fail
```

Since this transformation adds new constructors, all functions need to cover these cases, too. The new rules return `Fail` if the function's argument is a failed computation and distribute function calls to both branches if the argument is a choice.

One issue with this approach is that call-time choice is not implemented yet. If a choice is duplicated during evaluation, this information cannot be recovered later. Therefore, each `Choice` constructor has an additional ID argument that identifies the same choices. Since each choice needs a fresh ID, functions use an additional `IDSupply` argument when choices are created.

The evaluation of a non-deterministic value is implemented by transforming the value into a search tree which can be traversed with different search strategies. In the process, each choice ID's decision is stored and then repeated if the same ID is encountered again.

While this approach is useful when the host language supports laziness and sharing, another approach is necessary to model these effects when they are not built into the language.

Fischer et al. [2009] introduce a monadic representation of non-determinism that supports sharing and non-strict evaluation. Out of simplicity, the implementation idea is presented in Haskell, similar to the approach of the original authors, using the example of permutation sort. The algorithm consists of three components. Firstly, a function `insert` that inserts an element non-deterministically at every possible position within a list.

```
insert :: MonadPlus m => a -> [a] -> m [a]
insert x xs = return (x:xs)
    `mplus` case xs of
        []      -> mzero
        (y:ys) -> do zs <- insert x ys
                    return (y:zs)
```

The second part is the function `perm` that inserts the head of a given list into the permutations of the list's tail.

```
perm :: MonadPlus m => [a] -> m [a]
perm [] = return []
perm (x:xs) = do ys <- perm xs
                zs <- insert x ys
                return zs
```

Finally, the function `sort` generates permutations and then tests whether they are sorted.

```
sort :: MonadPlus m => [Int] -> m [Int]
sort xs = do ys <- perm xs
            guard (isSorted ys)
```

```
return ys
```

The function `isSorted` compares each element in a list to the next one to determine whether the list is sorted. When we test this implementation, we can see that the runtime increases significantly when adding even a few elements.

```
λ> sort [9, 8..1] :: [[Int]]
[[1,2,3,4,5,6,7,8,9]]
(0.69 secs)
λ> sort [10, 9..1] :: [[Int]]
[[1,2,3,4,5,6,7,8,9,10]]
(6.67 secs)
λ> sort [11, 10..1] :: [[Int]]
[[1,2,3,4,5,6,7,8,9,10,11]]
(77.54 secs)
```

The reason for the factorial runtime is that the implementation is needlessly strict. A list of length n has $n!$ permutations, all of which are generated when running `sort`. This matches our observation above, since adding a tenth element increases the runtime by a factor of 10 and an eleventh element multiplies the runtime of the ten-element list by eleven.

If we consider the implementation of `isSorted`, we can see that, as soon as the comparison of two elements yields `False`, the function returns `False` and does not evaluate the remainder of the list.

```
isSorted :: [Int] -> Bool
isSorted (x:y:zs) = (x <= y) && isSorted (y:zs)
isSorted _       = True
```

However, since we use `bind` to pass permutations from `perm` to `isSorted`, each permutation is fully evaluated before it is determined whether the permutation is sorted. This leads to the complete evaluation of every permutation, which results in an inefficient program.

Similarly, when we consider the Curry example `head (1 : head [] : [])`, the strictness of our `MonadPlus` approach shows again. The corresponding Haskell expression is as follows.

```
hd [] >>= \x -> hd (1 : x : [])
```

Here `hd :: MonadPlus a => [a] -> m a` is the lifted head function. Evaluating the expression in Haskell yields `mzero`, that is, no result, while Curry returns 1. The reason is the definition of the bind operator. For example, the monad instance for lists defines `bind` as `xs >>= f = concatMap f xs`. In the expression above, this means that the pattern matching within `concatMap` evaluates `hd []` to `mzero` and thus returns `mzero`.

The strictness observed in both examples is the motivation for an alternative approach. The problem with the above implementations is that non-deterministic arguments of constructors need to be evaluated completely before the computation can continue. There-

fore, we would like to be able to use unevaluated, non-deterministic computations as arguments of constructors.

As mentioned before, we can implement this idea by adapting all data types so that they may contain non-deterministic components.

```
data List m a = Nil | Cons (m a) (m (List m a))
```

The list data type now has an additional argument `m` of type `* -> *` that represents a non-determinism monad. Instead of fixed constructors like `Choice`, the monad `m` determines the structure and evaluation strategy of the non-determinism effect. Two smart constructors `cons` and `nil` make handling the new list type more convenient.

```
nil :: Monad m => m (List m a)
nil = return Nil

cons :: Monad m => m a -> m (List m a) -> m (List m a)
cons x y = return (Cons x y)
```

Adapting the permutation sort functions to the lifted data type requires us to replace `[]` with `List m`. However, this is not sufficient because the list itself can be the result of a non-deterministic computation. Therefore, an additional `m` is wrapped around every occurrence of `List`.

```
insert' :: MonadPlus m => m a -> m (List m a) -> m (List m a)
insert' mx mxs = cons mx mxs
  `mplus` mxs >>= \xs -> case xs of
    Nil      -> mzero
    Cons my mys -> cons my (insert' mx mys)

perm' :: MonadPlus m => m (List m a) -> m (List m a)
perm' ml = ml >>= \l ->
  case l of
    Nil -> nil
    Cons mx mxs -> insert' mx (perm' mxs)
```

Whenever pattern matching occurred in the original definition, we now use `bind` to extract a `List` value. Since this only evaluates flat non-determinism and not non-determinism that occurs in the components, non-strictness is upheld as much as possible.

All functions now take arguments of the same type they return. Thus, the definition of `sort` does not need `bind` in order to pass permutations to `isSorted`.

```
sort' :: MonadPlus m => m (List m Int) -> m (List m Int)
sort' xs = let ys = perm' xs in
  isSorted' ys >>= \sorted -> guard sorted >> ys
```

We are now able to take advantage of `isSorted`'s non-strict definition. The implementation generates permutations only if there is a chance that the permutation is sorted, that is, only recursive calls of `perm` that are demanded by `isSorted` are executed.

We reconsider the Curry example head `(1 : head [] : [])`. Since the `List` data type now takes monad values as arguments, we can write the example using the smart constructors and a lifted head function as follows.

```
λ> hd' (cons (return 1) (cons (hd' nil) nil))  
1
```

Because we do not need to use `bind` to get the result of `hd' nil`, the expression is not evaluated due to non-strictness and the result is equal to Curry's output.

Data types with non-deterministic components solve the problem of non-strictness because each component can be evaluated individually, instead of forcing the evaluation of the whole term. Unfortunately, this leads to a problem. When unevaluated components are shared via Haskell's built-in sharing, computations, rather than results, are being shared. This means that the results can be different each time the computation is evaluated, which contradicts the intuition of sharing.

The solution to this problem is an explicit sharing combinator `share :: m a -> m (m a)` that allows sharing the results of a computation in a non-strict way. Here, `m` is a `MonadPlus` instance, similar to the monad used in the definition of the data type, that supports sharing. Thus, `share` takes a computation and then returns a computation that returns the result, that is, the shared value. The reason for this nesting of monad layers is that, in short, the `share` combinator performs some actions that can be immediately executed by `bind` (the outer monad layer), while the inner monad layer should only be evaluated when needed. This is explained in more detail later. With the explicit sharing operator we can adapt `perm'` to share the generated permutations in order to achieve non-strictness in combination with sharing.

```
sort' :: MonadPlus m => m (List m Int) -> m (List m Int)  
sort' xs = do ys <- share (perm' xs)  
             sorted <- isSorted'  
             guard sorted  
             ys
```

The `share` operator must satisfy certain laws, which we discuss in section 4.4. The implementation of `share` is subject of the next chapter.

3 Call-Time Choice modelled in Haskell

Based on the ideas presented in the last chapter, we now want to model call-time choice, that is, non-strictness, sharing and non-determinism, in Haskell. We still use `MonadPlus` to parameterize our programs. However, instead of, for example, using the list instance to make non-determinism visible, we define an effect functor that can express many different effects, including non-determinism and sharing. This approach, as introduced by Wu et al. [2014], will also be the base of the Coq implementation shown in chapter 4.

TODO: Definition effect

For the implementation of call-time choice, we want to be able to express different effects within our programs. However, not every program contains effects. There are also pure programs that have no side-effects besides the computation of a value. A data type that represents such programs could look as follows.

```
data Void a = Return a
```

Here, `Void` means the absence of effects.

If we consider programs that contain effects, also called impure, like, non-determinism, a data type that represents such values could look like the following.

```
data ND a = Return a
          | Fail
          | Choice (ND a) (ND a)
```

This data type also has a constructor to model pure values but in addition, there are constructors that represent failed computations and the non-deterministic choice between two values. We could go on and list data types that model many more effects but the question is: Is it possible to create a data type that, if appropriately instantiated, behaves like the original effect functor? This would allow us to represent programs with many different effects using one compact data type.

Answering this question requires abstracting the concrete form of effect functors into a general program data type. As we saw in the examples above, we need a way to represent pure values in a program. Therefore, the first constructor of our new program data type should be `Return a` for the type `a`, that is, the result type of the program. To model effects like non-determinism, the program is parameterized over effect functors of type `* -> *` that represent, for example, `Fail` and `Choice`. We call this argument `sig` because the signature of a program tells us which effects can occur. So far, programs are defined as `data Prog sig a = Return a`. In order to make use of the `sig` component, we need to add a constructor for impure operations. The `ND` data type shows us that effect functors can be defined recursively. Thus, the constructor for impure programs should be recursive, too, to be able to represent this structure.

```
data Prog sig a = Return a | Op (sig (Prog sig a))
```

With this definition of `Prog`, we are able to represent the original functors by instantiating `sig` appropriately. For `Void`, we already have the `Return` constructor. Therefore, the data type we can use with `Prog` does not need a constructor anymore, that is, `data Void' a`.

```
VoidProg a = Return a
           | Op (Void' (VoidProg a)) -- Void' has no constructors!
```

The type `Prog Void'` now resembles the original type `Void` since the `Op` constructor would require a value of type `Void'`, which we cannot construct.¹ Only `Return` can be used to define values, similar to the original data type.

Similar to `Void'`, we can define a data type `Choice` that represents `Choice` in combination with `Prog`.

```
data ND p = Fail | Choice p p
```

Again, we can omit the `Return` constructor because it is already part of the `Prog` data type. For the same reason, the type variable `a` has been replaced with the variable `p`, since `ND` does not have values as arguments but rather programs that return values.

```
data NDProg a = Return a
              | Op (Choice (NDProg a))
```

Since `Op` applies `sig` recursively, this yields the following type, which is equivalent to the original data type.

```
data NDProg a = Return a
              | OpFail
              | OpChoice (NDProg a) (NDProg a)
```

We have found a way to model effect functors as instances of the data type `Prog`, which essentially models a tree with leafs, represented by the `Return` constructor, and branches that have the form defined by `sig`.

TODO: Tree structure visualization?

3.1 Free Monads

- What are free monads?
- Why do we use free monads?

The data type `Prog` is better known as the free monad. We saw in the previous chapter that `Free` can be used to model other data types. In addition, `Free` is a monad that can turn any functor into a monad.

¹It is possible to use `undefined` to create an impure value of type `Prog Void' a`. Since this is not possible in `Coq`, we do not consider this in the Haskell implementation.

We consider, for example, the type `Free One` where `data One a = One`. Here `a` is a phantom type that we need because `Free` expects a functor. The monad instance for `Free` is as follows.

```
data Free f a = Pure a | Impure (f (Free f a))

instance (Functor f) => Monad (Free f) where
  return = Pure
  Pure x >>= g = g x
  Impure fx >>= g = Impure (fmap (>>= g) fx)
```

Since `One` has only a single, non-recursive constructor `One`, the only possible impure value is `Impure One`, whereas the usual `Return` constructor remains. If `bind` encounters the value `One`, the function `g` is distributed deeper into the term structure using `fmap`. Since `fmap One = One`, it becomes apparent that the monad constructed by `Free One` is the `Maybe` monad.

Since we want to model different effects in our program, the free monad makes writing programs easier by allowing monadic definitions without defining a separate monad instance for each effect.

TODO: Visualization of `fmap` tree

TODO: Find more reasons for using free monads

TODO: Stacking monads does not necessarily yield monads again

TODO: Split `ND` into choice and fail instead of one

3.2 Modelling Effects

- Explanation of the `Prog/sig` infrastructure
- `ND` and state effect implementation

In the previous sections the free monad and its ability to represent effect functors was discussed. The goal of this section is to explore the infrastructure that allows us to combine multiple effects, write effectful programs and compute the result of such programs.

3.2.1 Combining Effects

Firstly, we would like to combine multiple effects. For this purpose, we use the technique introduced by Swierstra [2008] to define a data type that combines the effect functors `sig1` and `sig2`. The infix notation simplifies combining multiple effects via nested applications of `++:`.

```
data (sig1 :+: sig2) a = Inl (sig1 a) | Inr (sig2 a)
```

For example, the type `ND :+: One` is a functor that we can use with `Prog` to define programs that contain non-determinism and partiality as follows.

```
progNDOne :: Prog (ND :+: One) Int
progNDOne = Op (Inl (Choice (Op (Inr One)) (Return 42)))
```

In the example `progNDOne` we define a program that represents the non-deterministic choice between a program whose value is absent and a program that returns 42. The

complexity of nesting constructors of `Prog` and `:+:` correctly increases quickly for bigger terms. Therefore, we define a type class that allows us to define such expressions more conveniently. The class is parameterized over two functors, one of which is a subtype – regarding `:+:` – of the other.

```
class (Functor sub, Functor sup) => sub <: sup where
  inj :: sub a -> sup a
```

We need a few instances of the class `<:` to make it useful. The simplest case is `sig <: sig` where we want to inject a value of type `sig a` into the same type. Since we do not need to modify the value in any way, `id` is used to define `inj`.

```
instance Functor sig => sig <: sig where
  inj = id
```

The next instance covers the case `sig1 <: (sig1 :+: sig2)`. Since we already know that `sig1` is part of the sum type, we only need to apply the correct constructor of `:+:`, that is, `Inl` because `sig1` is the left argument.

```
instance (Functor sig1, Functor sig2) => sig1 <: (sig1 :+: sig2) where
  inj = Inl
```

The last instance assumes that we can inject `sig` into `sig2` and describes how we can inject `sig` into `sig1 :+: sig2`. In this case, we can use `inj` to receive a value of type `sig2 a`. All that remains is a situation similar to the previous instance, where we only need to use the matching constructor to complete the injection.

```
instance (Functor sig1, sig <: sig2) => sig <: (sig1 :+: sig2) where
  inj = Inr . inj
```

These instances allow us to write a polymorphic definition of the function `inject` which injects constructors depending on the given type of the program.

```
inject :: sig1 <: sig2 => sig1 (Prog sig2 a) -> Prog sig2 a
inject = Op . inj
```

`inject` can then be used as demonstrated in the following example.

```
λ> inject One :: Prog (One :+: ND) a
Op (Inl One)
λ> inject One :: Prog (ND :+: One) a
Op (Inr One)
```

The implementation of the function `inject` assumes that we can inject `sig1` into `sig2`. This is because `sig2` is the signature of the returned program and `sig1` is the type of the effect constructor that we want to inject. This restriction is justified because, for example, non-deterministic syntax should only appear in a program where `ND` is part of the signature. With this part of the infrastructure in place, we can redefine the example `progNDOne` without using `Inl` and `Inr` explicitly.

```

progNDOne' :: Prog (ND :+: One) Int
progNDOne' = inject (Choice (inject One) (Return 42))

```

Deriving the appropriate instance of `<:` when using `inject` is, however, not always unambiguous. The last two instances overlap in situations where `sig = sig1`. For example, the example `inject One :: Prog (One :+: One)` yields different values with respect to the chosen constructor of `<:`, depending on the instance.

```

λ> nothing
Op (Inl One) -- second instance
λ> nothing
Op (Inr One) -- third instance

```

This is because the type constraint of `inject`, in this case `One <: (One :+: One)`, matches both the second and third instance. Haskell does not accept overlapping instances by default, which is why we prioritize one instance via pragmas. In practice, the different term structure due to `Inl` and `Inr` does not influence the evaluation as long as we do not explicitly match for the constructors. This is ensured by an additional function `prj` of the type class `<:`, which is discussed in the next section.

3.2.2 Simplified Pattern Matching

While the function `inject` allows us to write programs in a more convenient way, we also need to consider how we can evaluate programs. The same issue of nested applications of `Op`, `Inl` and `Inr` applies when we want to distinguish different effects via pattern matching. Thus, we add a second function `prj` to the type class `<:`.

```

class (Functor sub, Functor sup) => sub <: sup where
  inj :: sub a -> sup a
  prj :: sup a -> Maybe (sub a)

```

The function `prj` is a partial inverse to `inj`. This means that we can project values of a type `sup a` into a subtype `sub a`. For this reason, the return type of the function is a `Maybe` type. Similar to `inj`, we have to define instances for the same cases as before.

- For `sig <: sig`, we can define `prj` as `Just` because we know that every element of the supertype is also an element the subtype.
- `sig1 <: (sig1 :+: sig2)` means that we can return `Just x` for `Inl x`. However, for `Inr` we need to return `Nothing` because we cannot, in general, project from `sig2` to `sig1`.
- In the last case `sig <: sig2 => sig <: (sig1 :+: sig2)` we know that we can project from `sig2` to `sig`. Thus, in case of `Inr x`, where `x` has the type `sig2`, we can apply `prj` to construct a value of appropriate type. The other case `prj (Inl _)` is handled by returning `Nothing`.

With the definition of `prj` and the instances of `:<:`, we can now define the function `project` which we can use to make pattern matching more convenient.

```
project :: (sub :<: sup) => Prog sup a -> Maybe (sub (Prog sup a))
project (Op s) = prj s
project _      = Nothing
```

Due to the recursive definition of the `Prog` data type, constructors like `Choice` have `Prog` arguments themselves. Thus, `sub` is applied to `Prog sup a` in the return type of the projection. We can only project effectful values because generally it is not clear which functor we should choose for `sub` when projecting a `Return` value.

Finally, we can now inject and project effectful values. Since `project` is a partial inverse of `inject`, the equation `project (inject x) = Just x` holds for values `x` of appropriate type, excluding failing computations. This is demonstrated in the following example.

TODO: Does it hold?

```
λ> type T = Maybe (ND (Prog (ND :+: One) Int))
λ> project (inject (Choice (Return 42) (Return 43))) :: T
Just (Choice (Return 42) (Return 43))
```

Now that we can use `project` as an abstraction of the concrete term structure regarding `:<:`, we can write a first function that evaluates a non-deterministic, partial program.

```
evalNDOne :: Prog (ND :+: One) a -> [a]
evalNDOne (Return x) = [x]
evalNDOne p = case project p of
    Just (Choice p1 p2) -> evalNDOne p1 ++ evalNDOne p2
    Just Fail            -> []
    Nothing              -> case project p of
        Just One -> []
        Nothing  -> []
```

When `evalNDOne` encounters a value `Return x`, `x` is returned as a singleton list. For effectful programs, we can use `project` to distinguish between the constructors of one effect at a time. The case patterns hold the necessary type information for `project`. When the projection returns `Nothing`, another effect can be matched in a nested case expression. Since we never need to explicitly match for `Inl` or `Inr`, overlapping patterns in the instances of `:<:` do not affect the evaluation of programs in our model.

Although we have already eliminated `Inl`, `Inr` and `Op` from functions that create or evaluate programs, there can be done even more to simplify programming with effects. Two language extensions, `PatternSynonyms` and `ViewPatterns`, allow us to write definitions like the following.

```
pattern PChoice p q <- (project -> Just (Choice p q))
```

View patterns – the right-hand side of the `<-`, make pattern-matching for certain cases more compact. A view pattern consists of a function on the left-hand side of `->`, that is applied to the value that the pattern is matched against, and a pattern on the right-hand

side. The result of the function call is matched against this pattern and the variables inside the pattern can be used in the definition. The function `evalNDOne` can be defined using view patterns in the following way.

```
evalNDOne' :: Prog (ND :+: One) a -> [a]
evalNDOne' (Return x) = [x]
evalNDOne' (project -> Just (Choice p1 p2)) = evalNDOne' p1 ++ evalNDOne' p2
evalNDOne' (project -> Just Fail             ) = []
evalNDOne' (project -> Just One              ) = []
```

We cannot use `(project -> Nothing)` without type annotations as a pattern because this would result in overlapping instances. However, no effects other than those specified in the signature can occur within the program. Therefore, the `Nothing` pattern is not necessary.

The second component of the pattern definition above is the option to define a synonym for more complex patterns. In this case, we name the view patterns similar to the original constructors of the effects. While this is necessary for every effect constructor, it allows us to rewrite the definition in the following way.

```
evalNDOne'' :: Prog (ND :+: One) a -> [a]
evalNDOne'' (Return      x) = [x]
evalNDOne'' (PChoice p q) = evalNDOne'' p ++ evalNDOne'' q
evalNDOne'' (PFail       ) = []
evalNDOne'' (POne        ) = []
```

Writing programs that evaluate effectful programs is now almost as convenient as simple pattern matching. Finally, a useful definition for working with programs that have the signature `f :+: g`, where we want to match for `f` but not `g`, is as follows.

```
pattern Other s = Op (Inr s)
```

Since `:+:` is right-associative in nested applications, we can match for the left argument effect and conveniently match all remaining effects with `Other`.

3.2.3 Handling Effects

For each effect in a program's signature, a handler is required. Handling an effect means transforming a program that contains a certain effect into a program where the effect's syntax does not occur anymore. However, the syntax is not just removed, but the effect's semantics is applied. The semantics of an effect is therefore given by its handler. In the following we discuss handlers for the effects non-determinism and state.

Void Effect We begin with the data type for absence of effects, `Void`. Due to its definition without constructors, there is no `Void` syntax that needs to be handled. The only constructor for programs with the signature `Void` is `Return`, which we can handle by returning the argument. Thus, the handler for `Void` removes the program layer and is usually applied last, when all other effects have been handled.

```
run :: Prog Void a -> a
run (Return x) = x
```

Non-determinism Effect We already defined a data type for non-deterministic programs in chapter 3. The Choice constructor did not contain any IDs, which we need for the implementation of call-time choice. Thus, the revised data type is as follows.

```
data ND p = Fail | Choice (Maybe ID) p p
```

Not every non-deterministic choice in a program needs an ID, since IDs slow down the evaluation of choices considerably. Thus, IDs are optional and only assigned when necessary, that is, when choices are shared.

In the last section, we already defined a function `evalNDOne` that handles the simple ND type without IDs by returning a list of results, where, for each choice, the result lists are concatenated. For choices with IDs, however, this is not sufficient. We begin by transforming the program into a program that returns a tree data type which mirrors the non-determinism structure.

TODO: Keep tree structure?

```
runND :: (Functor sig) => Prog (ND :+: sig) a -> Prog sig (Tree.Tree a)
runND (Return a) = return (Tree.Leaf a)
runND Fail      = return Tree.Failed
runND (Choice m p q) = do
  pt <- runND p
  qt <- runND q
  return (Tree.Choice m pt qt)
runND (Other op) = Op (fmap runND op)
```

Next, we need to memorize the decisions that were made while traversing the choice tree. For this reason, we define a data type `Decision` that indicates whether the left or right branch of a choice has been picked before for a particular choice ID. A `Memo` maps IDs to decisions.

```
data Decision = L | R
type Memo = Map.Map ID Decision
```

The depth-first traversal of the choice tree is implemented in the function `dfs`. The returned list of results is created similar to the approach in `evalNDOne`, except for the case where a choice has a non-empty ID. The ID could have appeared in a choice that is closer to the root node of the tree and thus, the choice could have already been decided. Therefore, we need to look up the ID in the `Memo`. If the choice has not been made yet, that is, `Nothing` is returned, the `Memo` is updated with `L` for the left branch and `R` for the right branch. The recursive calls then descend into the corresponding branch and will make the same decision for this ID if it occurs again. If, on the other hand, a decision is returned by the lookup function, the branch of the recursive call is chosen according to the decision.

```
dfs :: Memo -> Tree a -> [a]
```

```

dfs mem Failed = []
dfs mem (Leaf x) = [x]
dfs mem (Choice Nothing t1 t2) = dfs mem t1 ++ dfs mem t2
dfs mem (Choice (Just n) t1 t2) =
  case Map.lookup n mem of
    Nothing -> dfs (Map.insert n L mem) t1
              ++ dfs (Map.insert n R mem) t2
    Just L -> dfs mem t1
    Just R -> dfs mem t2

```

The function `dfs` is called with an empty map and yields the list of results that the choice tree represents.

TODO: Examples

State Effect Stateful computations are an important part of the sharing effect that is presented in section 3.3. We begin by defining the syntax of the state effect. Usually, stateful computations can read the current state with `get` and set a new state with `put`. Thus, the data type needs those two constructors, too. We add an additional type variable that abstracts the type of values that the state can hold. The variable `p` represents the program type as before.

```

data State s p = Get' -- ?
               | Put' -- ?

```

Both constructors have an effect on a program, that is, a scope in which the effects are visible. Thus, both constructors need a program argument `p`. For `Put'`, we can simply add the arguments `s` for the new state and `p` for the program in which the new state is set. The constructor `Get'`, however, contains the program in a different form, namely a function `s -> p`. The reason for this is that, if we were using a simple `p` argument, the handler would have to somehow replace all `get`-occurrences of the state with appropriate values. This would require evaluating the whole program, which would defeat the purpose of preserving non-strictness. Hence, the program is added to `get` in the form of a functional expression where the function argument replaces the occurrences of the state that are being read in the program. The data type for the state effect now looks as follows.

```

data State s p = Get' (s -> p)
               | Put' s p

```

The smart constructors for stateful programs are defined by instantiating the program and function arguments appropriately. For `get`, this means that we need to supply a function of type `s -> Prog sig s` since `p` is `Prog sig s` in this context. Conveniently, the return function matches this type and thus, is the initial argument of `Get'`. For `put`, the new state and a program that returns `()` are supplied to `Put'` because `put` does not return any information.

```

get :: (State s <: sig) => Prog sig s
get = inject (Get' return)

```

```
put :: (State s <: sig) => s -> Prog sig ()
put s = inject (Put' s (return ()))
```

The choice of initial function arguments might not seem intuitive at first because it is not clear how the remaining program finds its way into the argument of, for example, `Get'`. Therefore, we consider an example of the state effect and how the free monad is used to write programs.

```
p :: Prog (State Int :+: Void) Int
p = do put 42
      i <- get
      return (i * 2)
```

The program sets a state 42, gets the value of the current state and then returns double of that. The normal form of `p` can be computed by evaluating the occurrences of `bind`. We recall the monad instance for the free monad: `bind` uses `fmap` to distribute a function deeper into a term. Thus, we first define a `Functor` instance.

```
instance (Functor sig) => Monad (Prog sig) where
  return x = Return x
  Return x >>= f = f x
  Op op >>= f = Op (fmap (>>= f) op)

instance Functor (State s) where
  fmap f (Get' g) = Get' (f . g)
  fmap f (Put' s p) = Put' s (f p)
```

In the case of `Get'`, we need to apply `g` to a state in order to obtain a program that we can apply `f` to. Thus, we pass the result from `f` to `g` via function composition. For `Put'`, the state `s` remains unmodified and the function `f` is applied to the program argument `p` of the constructor.

Now we can transform the program `p` into normal form as follows.

```
put 42 >>= \_ -> get >>= \i -> return (i * 2)
= inject $ fmap (>>= \_ -> get >>= \i -> return (i * 2)) (Put' 42 (return ()))
= inject $ Put' 42 (get >>= \i -> return (i * 2))
= inject $ Put' 42 (inject $ fmap (>>= \i -> return (i * 2)) (Get' return))
= inject $ Put' 42 (inject $ (Get' (\i -> return (i * 2))))
```

`Op` as well as `Inl` and `Inr` constructors are replaced by `inject` in this example. The expression is transformed by applying the definitions of `bind` and `fmap`. In the last step, we simplify the expression by applying the left identity monad law, that is, $(\>>= f) \cdot \text{return} = f$.

We can now see that the remaining program after `get`, that is, the `return` call, has been moved into the argument function of `Get'`. The function expects a state and replaces the variables, that were bound to the return value of `get` in the original program, with the state.

Now that we have seen the definition of stateful program syntax and how the state flows through the program via functions, we can define the handler for the state effect.

Naturally, the handler needs to keep track of the current state, which is the first argument of the function. Then, the function expects a program that contains state syntax. Finally, the return type is a program that returns a pair of the current state and a return value.

```
runState :: Functor sig => s -> Prog (State s :+: sig) a -> Prog sig (s, a)
runState s (Return a) = return (s, a)
runState s (Get k) = runState s (k s)
runState s (Put s' k) = runState s' k
runState s (Other op) = Op (fmap (runState s) op)
```

For pure values, the current state and the value inside the Return constructor is returned. When a Get is encountered, we apply the function argument, which expects a state, to the current state and do a recursive call with the resulting program. Put is handled by a recursive call where the old state is replaced by the new state while the program stays the same. Finally, other syntax is handled by using fmap to distribute the handler deeper into the term structure, similar to the other handlers we have seen.

The example program p can now be handled by first calling the handler runState to handle the state effect, followed by run to extract the result from the program structure.

```
λ> run . runState 1 $ p
(42,84)
```

As expected, the first component represents the current state, which was set by put to 42, while the second component is the result that was returned after multiplying the current state by two.

Handling Order

3.3 Sharing

- How can we implement simple sharing as an effect?
- What about deep/nested sharing?
- Examples (exRecList, ...)

4 Call-Time Choice modelled in Coq

The goal of this chapter is to transfer the Haskell implementation of call-time choice to Coq. We begin with the data structure `Prog`, that is, the free monad, which allowed us to model programs with effects of type `sig` and results of type `a`.

```
data Prog sig a = Return a | Op (sig (Prog sig a))
```

The definition in Coq looks very similar to the Haskell version, aside from renaming and the explicit constructor types.

```
Inductive Free F A :=  
| pure : A -> Free F A  
| impure : F (Free F A) -> Free F A.
```

However, the definition is rejected by Coq upon loading the file with the following error message.

```
Non-strictly positive occurrence of "Free" in "F (Free F A) -> Free F A".
```

The reason for this error is explained in the next section.

4.1 Non-strictly Positive Occurrence

- What does non-strictly positive occurrence mean?
- Motivation for usage of containers

In section 2.1, we learned that Coq distinguishes between non-recursive definitions and functions that use recursion. The reason for this is that Coq checks functions for termination, which is an important part of Coq's proof logic. To understand why functions must always terminate in Coq, we consider the following function.

```
Fail Fixpoint loop (x : unit) : A := loop x.
```

The function receives an argument `x` and calls itself with the same argument. Since this function obviously never terminates, the result type `A` is arbitrary. In particular, we could instantiate `A` with `False`, the false proposition. The value `loop tt : False` could be used to prove anything, according to the principle of explosion. For this reason, Coq requires all recursive functions to terminate provably.

Returning to the original data type, what is link between `Free` and termination? It is well known that recursion can be implemented in languages without explicit recursion

syntax by means of constructs like the Y combinator or the data type Mu for type-level recursion.

```
Fail Inductive Mu A := mu : (Mu A -> A) -> Mu A.
```

Mu is not accepted by Coq for the same reason as Free: non-strictly positive occurrence of the respective data type. The problematic property of non-strictly positive data types is that the type occurs on the left-hand side of a constructor argument's function type. This would allow general recursion and thus, as described above, make Coq's logic inconsistent.

In case of Free, the non-strictly positive occurrence is not as apparent as before because the constructors do not have functional arguments. However, F is being applied to Free F A. If F has a functional argument with appropriate types, the resulting type becomes non-strictly positive, as shown below.

```
Definition Cont R A := (A -> R) -> R.
```

```
(* Free (Cont R) *)
Fail Inductive ContF R A :=
| pureC    : A -> ContF R A
| impureC  : ((ContF R A -> R) -> R) -> ContF R A.
```

In the type of impureC contains a non-strictly positive occurrence of ContF R A. Consequently, Coq rejects Free because it is not guaranteed that no instance violates the strict positivity requirement. Representing the Free data type therefore requires a way to restrict the definition to strictly positive data types. One approach to achieve this goal is described in the next section.

4.2 Containers

- How do containers work?
- How do we translate effect functors into containers?

Containers are an abstraction of data types that store values, with the property that only strictly positive data types can be modelled as a container. This will allow us to define a version of Free that works with containers of type constructors instead of the type constructors itself. First, however, we will have a more detailed look at containers.

The first component of a container is the type Shape. A shape determines how the data type is structured, regardless of the stored values. For example, the shape of a list is the same as the shape of Peano numbers: a number that represents the length of the list, or rather the number of Cons/Succ applications. A pair, on the other hand, has only a single shape.

The second component of a container is a function Pos : Shape -> Type that gives each shape a type that represents the positions within the shape. In the example of pairs, the shape has two positions, the first and second component. Each element of a list is a

position within the shape. Therefore, the position type for lists with length n is natural numbers smaller than n . Peano numbers do not have elements and therefore, the position type for each shape is empty.

Containers can be extended by a function that maps all valid positions to values. Since the position type depends on a concrete shape, the definition in Coq is quantified universally over values of type Shape.

```
Inductive Ext Shape (Pos : Shape -> Type) A :=  
  ext : forall s, (Pos s -> A) -> Ext Shape Pos A.
```

The extension of a container models the concrete data type.

4.3 Modelling Effects

- In which ways is the Coq implementation simplified, compared to Haskell?
- How does the adapted Prog/sig infrastructure work?
- How do we translate recursive functions?

4.4 Sharing

- Laws of sharing

5 Curry Programs modelled in Coq

- Can we use the Coq model of call-time choice to prove properties about actual Curry programs?

6 Conclusion

```
Class Monad (M: Type → Type) :=
{
  ret : ∀ A, A → M A;
  bind : ∀ A B, M A → (A → M B) → M B;
  left_unit : ∀ A B (x0: A) (f: A → M B),
    bind (ret x0) f = f x0;
  right_unit : ∀ A (ma: M A), bind ma (@ret A) = ma;
  bind_assoc: ∀ A B C (ma : M A) (f: A → M B) (g: B → M C),
    bind ma (fun y bind (f y) g) = bind (bind ma f) g
}.

Definition join (M: Type → Type) `(Monad M) A (mmx : M (M A)) : M A := bind _

End MonadClass.
Arguments join { _ } { _ } { _ }.

Section MonadInstance.

Variable F : Type → Type.
Variable C__F : Container F.
```

Bibliography

Bernd Braßel, Michael Hanus, Björn Peemöller, and Fabian Reck. KiCS2: A new compiler from curry to haskell. In Proceedings of the 20th International Conference on Functional and Constraint Logic Programming, WFLP'11, pages 1–18, Berlin, Heidelberg, 2011. Springer-Verlag. ISBN 978-3-642-22530-7. URL <http://dl.acm.org/citation.cfm?id=2032603.2032605>.

Sebastian Fischer, Oleg Kiselyov, and Chung-chieh Shan. Purely functional lazy non-deterministic programming. SIGPLAN Not., 44(9):11–22, August 2009. ISSN 0362-1340. doi: 10.1145/1631687.1596556. URL <http://doi.acm.org/10.1145/1631687.1596556>.

Wouter Swierstra. Data types à la carte. J. Funct. Program., 18(4):423–436, July 2008. ISSN 0956-7968. doi: 10.1017/S0956796808006758. URL <http://dx.doi.org/10.1017/S0956796808006758>.

Nicolas Wu, Tom Schrijvers, and Ralf Hinze. Effect handlers in scope. ACM SIGPLAN Notices, 49, 09 2014. doi: 10.1145/2633357.2633358.